

Campuswire

https://campuswire.com/c/GC36E5829/feed/87

Golam Mahmud Samdani 283 2 days ago asked a question

Visible to: Everyone

Resolved

MD Model Prediction #87

I'm trying to understand how molecular dynamics simulations are used to train machine learning models for predicting interatomic potentials—how do we ensure the model captures long-range interactions and remains accurate across different phases or defect structures?

1 upvotes, 0 comments, 77 views, 29 likes

Answers **Comments** **Answer this question**

Student answers **Best**

An instructor (Andre Schleife) endorsed this answer

Jahid Emon 1524 answered this question 2 days ago

To ensure machine-learned interatomic potentials trained on molecular dynamics data capture long-range interactions and remain accurate across phases or defects, researchers often include diverse training data that spans multiple configurations—equilibrium, strained, defected, and various phases. Models like MACE or NeqIP can incorporate long-range physics through specialized architectures or explicit electrostatics. Additionally, techniques like active learning help adaptively sample regions where the model is uncertain, improving generalization to unseen structures.

Reply

An instructor (Andre Schleife) endorsed this answer

Wenwen LENG 59 answered this question 2 days ago

Traditional MLIPs (e.g., standard Deep Potential, MTP, GAP) typically rely on **local descriptors**, with only finite cutoff radius (e.g., 6–10 Å). I know that DeePMD group proposed an extension to their Deep Potential model by incorporating long-range electrostatics using a physically inspired approximation:

The long-range Coulomb interaction between ions and valence electrons is modeled as interactions between **Gaussian charge distributions** placed at atomic and electronic centers.

[arXiv:2112.13327](#)

This allows the model to:

- retain local symmetry-preserving neural networks,
- while adding an electrostatic baseline with distance-dependent interactions,
- and handle **heterogeneous charge distributions** more realistically.

In practice, if long-range interactions are important for your system, you'll need to:

- Use or extend a model that includes explicit electrostatics (e.g., DP-ZBL, DP-LR, PhysNet),
- Or combine MLIPs with **charge equilibration schemes** (QE, HIPNN-QM9, etc.),
- And ensure your training set includes **diverse environments** (e.g., defects, interfaces, multiple phases) to promote generalizability.

Hope that helps.

1 reply

Golam Mahmud Samdani 283 2 days ago

Thanks for explaining this nicely!

Wenwen Active

1 online now

Campuswire

https://campuswire.com/c/GC36E5829/feed/64

Class feed

MSE598ML: Machine Learning for MatSE

My posts New post

Last week

How Transferable Are Equilibrium-Trained ML Potentials to Nonequilibrium Regimes? #64

Wenwen LENG 59 🐣 asked a question 9 days ago

Visible to: Everyone

Resolved

General

Hi everyone,

Another question for me. We often evaluate machine-learned interatomic potentials using benchmarks like **OC20** and **Matbench**. These datasets are widely adopted, but they're dominated by **relaxed or near-equilibrium structures**.

This leads to a question that's been bothering me for a while:

Can MLIPs trained (or pre-trained) on equilibrium data be trusted in nonequilibrium scenarios?

Examples include:

- highly strained structures,
- high-energy defects or collisions,
- cascade events or fracture simulations.

Even for state-of-the-art models like **MACE**, **NequiP**, or **GemNet**, if the training distribution doesn't cover these regions, aren't we just *extrapolating blindly*?

I'm aware of two typical strategies to improve model performance in such cases:

- ✓ Active learning / on-the-fly sampling
- ✓ Fine-tuning with domain-specific high-energy data

But both essentially boil down to *adding more data*.

So here's my real question:

If we always need to add new data for each task, how "general" are general-purpose MLIPs in practice?

Are pretrained models truly transferable, or just good warm starts?

Thanks!

0 0 82 31

Answers Comments

Student answers Best

An instructor (Andre Schleife) endorsed this answer

Jahid Emon 1524 🐣 answered this question 8 days ago

That's a great question, and honestly, a lot of people working with MLIPs are wrestling with the same thing. You're right – most benchmark datasets like OC20 or Matbench are heavily skewed toward relaxed or near-equilibrium structures, which means MLIPs trained on them can struggle when pushed into high-strain, high-energy, or defect-heavy regimes. Even top models like MACE or NequiP will likely be extrapolating if those regions weren't in the training set. Active learning and fine-tuning definitely help, but as you said, they're just smarter ways of adding more data. So in practice, these "general-purpose" MLIPs often end up being more like solid starting points that still need tailoring for specific applications.

Transferability sounds great in theory, but in high-stakes simulations like fracture or cascades, it's risky to rely on pretrained models alone without domain-specific validation.

1 reply Reply

Wenwen LENG 59 🐣 8 days ago

Thanks so much for your thoughtful response – that really clarified a lot!

As a follow-up: if I fine-tune or distill a general-purpose MLIP using my domain-specific data, I'd expect it to require **less training data** than training a model from scratch.

But in that case, wouldn't the resulting potential (after distillation) still be **structurally more complex** than a model trained only on my own dataset (especially when used in MD simulations), since it inherits a more unified and general representation of atomic environments?

Just wondering – is that added complexity a trade-off worth considering in practice?

To be clear, I'm not against distillation or transfer learning – I'm just trying to understand how much **computational cost or inference overhead** we're introducing when we use a distilled "general" model in production simulations.

Reply

3 online now

Rooms

Files

Grades

Settings

Collapse

Wenwen Active

Campuswire

https://campuswire.com/c/GC36E5829/feed/63

Class feed

MSE598ML: Machine Learning for MatSE

Notifications 10

DMS

Calendar

Search

Class feed

Rooms

Files

A+ Grades

Settings

Collapse

Wenwen Active

Wenwen LENG 59 🐣 asked a question 9 days ago

Visible to: Everyone

Resolved

Discussion: Single vs Double Precision in ML-based Interatomic Potentials #63

General

Hi everyone,

I'd like to share an observation and raise a discussion regarding the use of **single (float32) vs double precision (float64)** in machine learning-based interatomic potentials – particularly in **DeepMD-kit**, though I suspect this issue may be more general.

Observation

During the **training phase** of DeepMD neural network models, we typically observe **little to no performance difference** between using `float64` and `float32`.

However, once the trained model is **deployed** – for example, in **LAMMPS** for production-scale molecular dynamics simulations – the **performance gap becomes substantial**, often with `float32` being significantly faster (e.g., 2x–8x speedup).

Questions I'd like to discuss:

- Why is `float32` significantly faster than `float64`, especially on GPUs?
- Why does this precision difference appear negligible during training, but become prominent during inference?
- Is this performance difference specific to DeepMD-kit, or common to all **deep neural network-based MLIPs**, such as **MACE**, **NequIP**, or **SchNet**?
- Is this difference inherent to NN-based inference, or can it be mitigated through tools or model conversion?
- Do non-deep-learning ML models (e.g., moment tensor-based **MTP**, Gaussian process-based **GAP/GPUMD**) also suffer from performance or stability issues related to precision choice?

If you've encountered this issue or have experience converting models for lower-precision inference, I'd love to hear how you handled it.

What's your preferred trade-off between **accuracy** and **speed** in production-level MLIP simulations?

Thanks!

1 upvote 0 comments 107 views 32 people

Answers Comments Answer this question

Student answers

Jahid Emon 1524 🌏 answered this question 8 days ago

Float32 is significantly faster than float64, especially on GPUs, because modern deep learning hardware (like NVIDIA Tensor Cores) is highly optimized for single precision. During training, the precision difference is less noticeable due to the dominance of complex operations and mixed-precision strategies, but during inference, float32's speed advantage becomes clear. This behavior isn't unique to DeepMD-kit—it also applies to other MLIP frameworks like MACE or NequIP. In contrast, classical ML models like GAP or MTP are often CPU-bound and use float64 more for stability than speed. A common practice is to train in float64 for accuracy, then convert the model to float32 for faster, large-scale simulations when numerical stability allows.

1 reply Reply

Wenwen LENG 59 🐣 8 days ago

Thanks a lot for your explanation!

Just to clarify – in your view, does **double precision** typically provide **better accuracy during training**?

Is that why you mentioned that "*a common practice is to train in float64 and convert to float32 for inference*"?

If so, it makes sense on why not just use single precision from the start on training.

1 reply Reply

Jahid Emon 1524 🌏 8 days ago

in float32, the max of the bit is 128, and in float64 have 256, if I am not wrong. if you start with float32 there is a chance you get lower bound, e.g. 5.3258 if you use float64 you get 5.326 but if you use float32 then you get 5.325, as max bit size exceeds so truncated not round off.

1 reply Reply

Wenwen LENG 59 🐣 8 days ago

great

+2 online now